

Comprehensive Analysis of Min-Hashing and LSH Implementation

Sri Ganesh Thota (B22CS054)

Feb 14, 2025

Technical Report

1. K-Gram Generation and Jaccard Similarity (KGramJaccard.java)

Implementation Overview

The implementation addresses the requirement to create three types of k-grams: - Character-based 2-grams - Character-based 3-grams - Word-based 2-grams

Technical Implementation

```
// Character k-grams generation
public static Set<String> generateKGrams(String text, int k) {
    Set<String> kGrams = new HashSet<>();
    for (int i = 0; i <= text.length() - k; i++) {
        kGrams.add(text.substring(i, i + k));
    }
    return kGrams;
}

// Word k-grams generation
private static Set<String> generateWordKGrams(String text, int k) {
    String[] words = text.split("\\s+");
    Set<String> kGrams = new HashSet<>();
    for (int i = 0; i <= words.length - k; i++) {
        kGrams.add(String.join(" ", Arrays.copyOfRange(words, i, i +
k)));
    }
    return kGrams;
}
```

Key features: - HashSet ensures unique k-grams (addressing duplicate elimination requirement) - Handles 27-character alphabet (lowercase letters and space) - Efficient string processing using substring and split operations

2. Min-Hashing Implementation (MinHash.java)

Task Requirements

Build min-hash signatures using $t = \{20, 60, 150, 300, 600\}$ hash functions for D1 and D2 comparison.

Technical Implementation

```
private static int[][] computeMinHashSignatures(List<Set<String>>
sets, int t) {
    int[][] signatures = new int[sets.size()][t];
    Random rand = new Random();

    // Generate t random hash functions
    for (int i = 0; i < t; i++) {
        int a = rand.nextInt(M - 1) + 1;
        int b = rand.nextInt(M);
        // Compute minimum hash values
        for (int j = 0; j < sets.size(); j++) {
            signatures[j][i] = computeMinHash(sets.get(j), a, b);
        }
    }
    return signatures;
}
```

3. LSH Implementation Analysis (LSH.java)

Configuration

- Total hash functions: 160
- Target similarity threshold (τ): 0.7
- Parameters:
 - Bands (b): 8
 - Rows per band (r): 20

S-Curve Analysis

The implementation uses the formula:

$$f(s) = 1 - (1 - s^r)^b$$

where: - s is the actual Jaccard similarity - r is rows per band - b is number of bands

4. MovieLens Min-Hashing Results

Analysis of 5 runs with different numbers of hash functions:

50 Hash Functions:

Average False Positives: 73.84 pairs
Average False Negatives: 0.48 pairs

100 Hash Functions:

Average False Positives: 23.92 pairs
Average False Negatives: 0.32 pairs

200 Hash Functions:

Average False Positives: 23.32 pairs
Average False Negatives: 0.12 pairs

Performance Analysis

1. Accuracy Improvement:
 - 50 → 100 hash functions: 67.6% reduction in false positives
 - 100 → 200 hash functions: 2.5% reduction in false positives
 - False negatives consistently low across all configurations
2. Optimal Configuration:
 - 100 hash functions provides the best balance between accuracy and computational cost
 - Diminishing returns observed when increasing to 200 hash functions
 - 50 hash functions shows too many false positives to be reliable

5. LSH on MovieLens Dataset

Configuration Results

1. 50 Hash Functions (r=5, b=10):

Average Results for similarity ≥ 0.6 :
- False Positives: 42.8
- False Negatives: 0.4

2. 100 Hash Functions (r=5, b=20):

Average Results for similarity ≥ 0.6 :
- False Positives: 26.0
- False Negatives: 0.2

3. 200 Hash Functions (r=5, b=40):

Average Results for similarity ≥ 0.6 :
- False Positives: 45.8
- False Negatives: 0.2

4. 200 Hash Functions (r=10, b=20):

Average Results for similarity ≥ 0.6 :
- False Positives: 22.2
- False Negatives: 0.0

Impact of Higher Similarity Threshold (0.8)

The results show significant improvements when increasing the similarity threshold to 0.8:

1. 50 Hash Functions (r=5, b=10):

Average Results:

- False Positives: 73.2
- False Negatives: 0.4

2. 100 Hash Functions (r=5, b=20):

Average Results:

- False Positives: 25.0
- False Negatives: 0.6

3. 200 Hash Functions (r=10, b=20):

Average Results:

- False Positives: 14.0
- False Negatives: 0.2

Conclusions and Recommendations

1. Optimal Hash Function Count:

- 100 hash functions provides the best balance between accuracy and computational cost
- Using 200 hash functions shows minimal improvement in accuracy while doubling computation time

2. LSH Configuration:

- For similarity threshold 0.6:
 - Best configuration: r=10, b=20 with 200 hash functions
 - Achieves lowest false negative rate (0.0) with reasonable false positives
- For similarity threshold 0.8:
 - Performance improves across all configurations
 - False positives decrease significantly
 - Slight increase in false negatives, but still within acceptable range

3. Trade-offs:

- More hash functions → Better accuracy but higher computation cost
- More bands → Higher probability of finding similar pairs but more false positives
- Higher similarity threshold → Better precision but might miss some similar pairs

4. Implementation Efficiency:

- Use of HashSet for k-gram storage ensures O(1) lookup
- Efficient signature computation using linear hash functions
- Effective band-based LSH implementation reduces comparison space

Overall Conclusions

Task 1:

K-gram analysis reveals that character-based grams are effective in capturing overlap between similar documents, whereas word-based grams highlight semantic distinctions. This insight helps in selecting an

appropriate similarity measure based on specific application requirements.

Task 2:

Min-Hashing offers consistent similarity approximations, even when using a moderate number of hash functions. Increasing the number of hashes enhances accuracy but comes at the cost of higher computational complexity.

Task 3:

Fine-tuning hyperparameters in LSH is essential for optimal performance. By analyzing the derivative at $\tau = 0.7$, we identified $(r, b) = (20, 8)$ as an effective configuration for distinguishing similar from dissimilar document pairs. The candidate probability function supports that only highly similar documents (e.g., D1 and D2) are frequently selected.

Tasks 4 & 5:

Experiments on the MovieLens dataset demonstrate that increasing the similarity threshold (t) reduces false positives, while strategic banding (adjusting r and b) improves candidate selection. The LSH results suggest that while false positives can be nearly eliminated, false negatives depend on both the similarity threshold and the chosen banding strategy

These results demonstrate successful implementation of both Min-Hashing and LSH techniques, with clear trade-offs between accuracy and computational efficiency across different configurations.

Note:

- Feel Free to refer to `out.txt`, `out1.txt` and `filtered_output.txt` for more information about outputs achieved by my code.
- `read.py` was used to filter the output from `out1.txt` to `filtered_output.txt`.